

```
import os
import shutil
import zipfile
from sklearn.model_selection import train_test_split
import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras.preprocessing.image import ImageDataGenerator, load_img, img_to_array
from tensorflow.keras.applications import ResNet50
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Dense, Dropout, GlobalAveragePooling2D
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping
import tensorflow as tf
```

```
#Download dataset using Kaggle CLI
!mkdir -p ~/.kaggle
!echo '{"username":"","key":""}' > ~/.kaggle/kaggle.json
!chmod 600 ~/.kaggle/kaggle.json
!kaggle datasets download -d omkargurav/face-mask-dataset
```

```
# Unzip
with zipfile.ZipFile("face-mask-dataset.zip", 'r') as zip_ref:
    zip_ref.extractall("face_mask_data")
```

 Dataset URL: <https://www.kaggle.com/datasets/omkargurav/face-mask-dataset>
License(s): unknown

```
# Prepare dataset
source_dir = '/content/face_mask_data/data'
base_dir = '/content/face_mask_data/split_data'
train_dir = os.path.join(base_dir, 'train')
val_dir = os.path.join(base_dir, 'val')

for category in ['with_mask', 'without_mask']:
    os.makedirs(os.path.join(train_dir, category), exist_ok=True)
    os.makedirs(os.path.join(val_dir, category), exist_ok=True)

    img_dir = os.path.join(source_dir, category)
    images = os.listdir(img_dir)
    train_imgs, val_imgs = train_test_split(images, test_size=0.2, random_state=42)

    for img in train_imgs:
        shutil.copy(os.path.join(img_dir, img), os.path.join(train_dir, category, img))
    for img in val_imgs:
        shutil.copy(os.path.join(img_dir, img), os.path.join(val_dir, category, img))
```


```
# Image settings
img_size = 224
batch_size = 32

# Data generators
train_datagen = ImageDataGenerator(rescale=1./255,
                                    rotation_range=20,
                                    zoom_range=0.2,
                                    horizontal_flip=True)

val_datagen = ImageDataGenerator(rescale=1./255)

train_gen = train_datagen.flow_from_directory(train_dir,
                                              target_size=(img_size, img_size),
                                              batch_size=batch_size,
                                              class_mode='binary')

val_gen = val_datagen.flow_from_directory(val_dir,
                                          target_size=(img_size, img_size),
                                          batch_size=batch_size,
                                          class_mode='binary')
```

 Found 6042 images belonging to 2 classes.
Found 1511 images belonging to 2 classes.

```
# Build model using ResNet50
base_model = ResNet50(input_shape=(img_size, img_size, 3),
                      include_top=False,
                      weights='imagenet')

base_model.trainable = False # Phase 1: freeze base

# Add classification head
x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dropout(0.3)(x)
x = Dense(128, activation='relu')(x)
x = Dropout(0.3)(x)
output = Dense(1, activation='sigmoid')(x)

model = Model(inputs=base_model.input, outputs=output)

model.compile(optimizer=Adam(learning_rate=0.001),
              loss='binary_crossentropy',
              metrics=['accuracy'])

# Train phase 1
model.fit(train_gen, validation_data=val_gen, epochs=5)
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet50_weights_tf_dim_ordering_tf_kernels_n94765736/94765736 0s 0us/step

/usr/local/lib/python3.11/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:121: UserWarning: Your `PyDataset` class self._warn_if_super_not_called()
Epoch 1/5
189/189 0s 445ms/step - accuracy: 0.5244 - loss: 0.7460/usr/local/lib/python3.11/dist-packages/PIL/Image.py:1043: L
warnings.warn(
189/189 110s 505ms/step - accuracy: 0.5245 - loss: 0.7458 - val_accuracy: 0.5963 - val_loss: 0.6656
Epoch 2/5
189/189 121s 437ms/step - accuracy: 0.5592 - loss: 0.6868 - val_accuracy: 0.6248 - val_loss: 0.6543
Epoch 3/5
189/189 83s 440ms/step - accuracy: 0.5661 - loss: 0.6789 - val_accuracy: 0.6115 - val_loss: 0.6540
Epoch 4/5
189/189 81s 430ms/step - accuracy: 0.5980 - loss: 0.6654 - val_accuracy: 0.6334 - val_loss: 0.6441
Epoch 5/5
189/189 83s 442ms/step - accuracy: 0.6051 - loss: 0.6581 - val_accuracy: 0.6267 - val_loss: 0.6443
<keras.src.callbacks.history.History at 0x79d5ee61b350>

```
# Phase 2: fine-tuning
base_model.trainable = True
for layer in base_model.layers[:100]: # freeze early layers
    layer.trainable = False

model.compile(optimizer=Adam(learning_rate=0.0001),
              loss='binary_crossentropy',
              metrics=['accuracy'])

model.fit(train_gen, validation_data=val_gen, epochs=10)
```

Epoch 1/10
189/189 138s 551ms/step - accuracy: 0.7070 - loss: 0.5732 - val_accuracy: 0.7697 - val_loss: 0.4684
Epoch 2/10
189/189 87s 461ms/step - accuracy: 0.8500 - loss: 0.3495 - val_accuracy: 0.6671 - val_loss: 0.8963
Epoch 3/10
189/189 87s 463ms/step - accuracy: 0.8570 - loss: 0.3171 - val_accuracy: 0.5917 - val_loss: 1.3428
Epoch 4/10
189/189 88s 466ms/step - accuracy: 0.9081 - loss: 0.2354 - val_accuracy: 0.8882 - val_loss: 0.2652
Epoch 5/10
189/189 141s 460ms/step - accuracy: 0.9168 - loss: 0.2144 - val_accuracy: 0.6208 - val_loss: 1.5184
Epoch 6/10
189/189 87s 458ms/step - accuracy: 0.9292 - loss: 0.1800 - val_accuracy: 0.8875 - val_loss: 0.2906
Epoch 7/10
189/189 87s 458ms/step - accuracy: 0.9393 - loss: 0.1607 - val_accuracy: 0.9199 - val_loss: 0.2011
Epoch 8/10
189/189 87s 459ms/step - accuracy: 0.9326 - loss: 0.1595 - val_accuracy: 0.5612 - val_loss: 2.5810
Epoch 9/10
189/189 89s 468ms/step - accuracy: 0.9505 - loss: 0.1323 - val_accuracy: 0.9431 - val_loss: 0.1480
Epoch 10/10
189/189 86s 455ms/step - accuracy: 0.9448 - loss: 0.1464 - val_accuracy: 0.8709 - val_loss: 0.5431
<keras.src.callbacks.history.History at 0x79d5303f3a10>

```
# Evaluate
loss, acc = model.evaluate(val_gen)
print(f"Validation Accuracy: {acc: .2f}")
```

```
print('Validation Accuracy: {:.2f}%')
```

48/48 ————— 4s 86ms/step - accuracy: 0.8576 - loss: 0.5807
Validation Accuracy: 0.87

```
# Test with real image
def predict_image(path):
    img = load_img(path, target_size=(img_size, img_size))
    img_array = img_to_array(img) / 255.0
    img_array = np.expand_dims(img_array, axis=0)
    pred = model.predict(img_array)[0][0]
    label = "With Mask" if pred < 0.5 else "Without Mask"
    print(f"Prediction: {label} ({pred:.2f})")
    plt.imshow(img)
    plt.title(label)
    plt.axis('off')
    plt.show()
```

```
from google.colab import files
```

```
# Upload an image file
uploaded = files.upload()
```

Choose files No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

```
-----
TypeError                                 Traceback (most recent call last)
<ipython-input-1-0c19da86b94d> in <cell line: 0>()
      2
      3 # Upload an image file
----> 4 uploaded = files.upload()
      5
```

1 frames

```
/usr/local/lib/python3.11/dist-packages/google/colab/files.py in _upload_files(multiple)
    169 files = _collections.defaultdict(bytes)
    170
--> 171 while result['action'] != 'complete':
    172     result = _output.eval_js(
    173         'google.colab._files._uploadFilesContinue("{output_id}")'.format(
```

```
TypeError: 'NoneType' object is not subscriptable
```

```
# Get the filename
filename = list(uploaded.keys())[0]
```

```
# Call the prediction function
predict_image(filename)
```

1/1 ————— 0s 45ms/step
Prediction: Without Mask (0.99)

Without Mask



